

Agentic Workflow Architecture

Overview

Hybrid workflow system combining:

Agent Orchestration & SLA Management

Agent Orchestration

- Agents are orchestrated per workflow using configurable strategies:
 - Execution modes: parallel, sequential, conditional, priority-based.
 - Red flag handling: early termination, escalation, or enhanced review.
 - Enrichment: agent outputs enrich entity metadata for downstream rules.
- Orchestration strategy is set per workflow and can be dynamically adjusted.

SLA Management

- SLA (Service Level Agreement) is tracked for each approval step and agent task.
- SLA breaches trigger escalation, notifications, or auto-approval/rejection.
- SLA configuration includes:
 - Per-agent and per-approver timeouts.
 - Escalation rules (who, when, how).
 - Fallback actions (auto-escalate, auto-reject, notify).
- SLA status is visible in workflow state and audit logs.

Example: Agent Orchestration & SLA in Workflow Template (YAML)

```
workflow:
  version: 2.2
  agents:
    - agentId: "kyc-fraud-agent"
      type: "MCP"
      config:
        endpoint: "http://fraud-service:8090"
        model: "claude-3-7-sonnet"
        tools: ["analyze_fraud", "check_id"]
        timeout: 20000
      orchestration:
        mode: "ASYNC_RED_FLAG"
        priority: 1
        slaHours: 6
        escalation:
          after: "PT6H"
          escalateTo: "FRAUD_TEAM"
          notify: true
    - agentId: "document-ocr-agent"
```

```

    type: "CUSTOM"
    config:
      endpoint: "http://ocr-service:8096"
      tools: ["extract_text", "validate_format"]
      timeout: 10000
    orchestration:
      mode: "SEQUENTIAL"
      priority: 2
      slaHours: 2
      escalation:
        after: "PT2H"
        escalateTo: "OPS_TEAM"
        notify: true
  slaManagement:
    globalTimeout: "PT24H"
    breachAction: "ESCALATE"
    notifyOnBreach: true

```

Example: Agent Orchestration & SLA in Workflow Template (JSON)

```

{
  "templateId": "DOCUMENT_VERIFICATION_V2",
  "version": "2.2",
  "entityType": "DOCUMENT_VERIFICATION",
  "agentConfig": {
    "enableAgents": true,
    "agents": [
      {
        "agentId": "kyc-fraud-agent",
        "type": "MCP",
        "config": {
          "endpoint": "http://fraud-service:8090",
          "model": "claude-3-7-sonnet",
          "tools": ["analyze_fraud", "check_id"],
          "timeout": 20000
        },
        "orchestration": {
          "mode": "ASYNC_RED_FLAG",
          "priority": 1,
          "slaHours": 6,
          "escalation": {
            "after": "PT6H",
            "escalateTo": "FRAUD_TEAM",
            "notify": true
          }
        }
      },
      {
        "agentId": "document-ocr-agent",

```

```

    "type": "CUSTOM",
    "config": {
      "endpoint": "http://ocr-service:8096",
      "tools": ["extract_text", "validate_format"],
      "timeout": 10000
    },
    "orchestration": {
      "mode": "SEQUENTIAL",
      "priority": 2,
      "slaHours": 2,
      "escalation": {
        "after": "PT2H",
        "escalateTo": "OPS_TEAM",
        "notify": true
      }
    }
  },
  "slaManagement": {
    "globalTimeout": "PT24H",
    "breachAction": "ESCALATE",
    "notifyOnBreach": true
  }
},
"decisionTables": [ ... ]
}

```

Decision Flow Patterns

Pattern 1: Rule-Based DMN (Traditional)

Submit → Validate → Evaluate DMN Rules → Assign Approvers → Human Approval

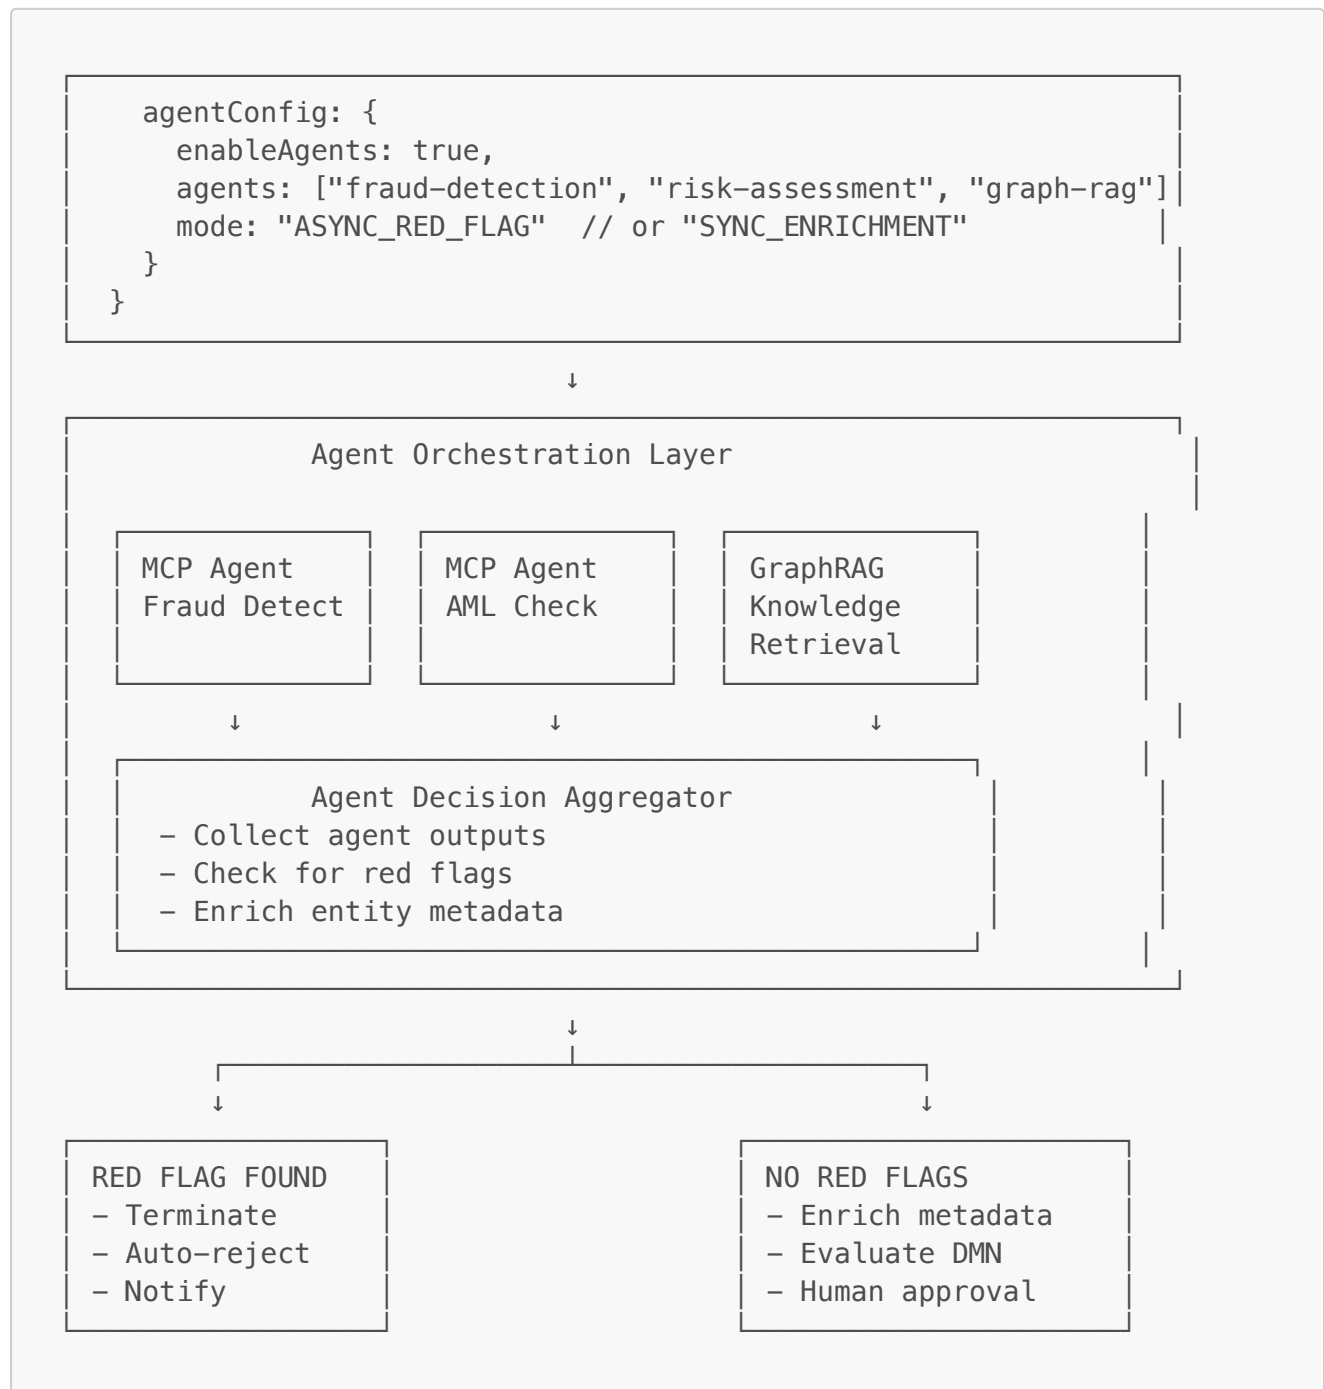
Pattern 2: Async Agent Red Flag Detection

Submit → Validate → Launch Agents (async) → Red Flag? → TERMINATE
 ↓
 No Red Flags → Continue to DMN → Human Approval

Pattern 3: Sync Agent + Rule Hybrid

Submit → Validate → Agent Analysis → Extract Insights → Evaluate DMN with Agent Data → Human Approval

Architecture



Agent Types

1. MCP Agent (Model Context Protocol)

Purpose: Real-time analysis using external AI services

Example: Fraud Detection Agent

```
{  
  "agentId": "fraud-detection-mcp",  
  "type": "MCP",  
}
```

```

"config": {
  "mcpServerUrl": "http://fraud-detection-service:8090",
  "timeout": 30000,
  "model": "claude-3-7-sonnet",
  "tools": [
    "analyze_transaction_patterns",
    "check_device_fingerprint",
    "verify_behavioral_biometrics"
  ]
},
"mode": "ASYNC_RED_FLAG",
"redFlagConditions": {
  "fraudProbability": "> 0.7",
  "suspiciousPatterns": "count > 3",
  "deviceMismatch": "true"
},
"enrichmentOutputs": [
  "fraudScore",
  "riskFactors",
]
}

```

2. GraphRAG Agent

```

```json
 "agentId": "compliance-graph-rag",
 "type": "GRAPH_RAG",
 "config": {
 "embeddingModel": "text-embedding-3-large",
 "retrievalDepth": 3,
 "regulations",
 "high_risk_countries",
]
 "mode": "SYNC_ENRICHMENT",
 {
 "cypher": "MATCH (c:Customer)-[:RELATED_TO]->(s:SanctionedEntity)
WHERE c.id = $customerId RETURN s"
 },
 {
 "name": "get_jurisdiction_rules",
 "cypher": "MATCH (j:Jurisdiction {country: $country})-
[:HAS_REGULATION]->(r:Regulation) RETURN r"
 }
],
"enrichmentOutputs": [
 "sanctionsMatch",
 "applicableRegulations",
 "complianceRequirements",
 "historicalRiskIndicators"
]
}

```

### 3. MCP Reasoning Agent

**Purpose:** Multi-step reasoning for complex decisions

**Example: Risk Assessment Agent**

```
{
 "agentId": "risk-assessment-mcp",
 "type": "MCP",
 "config": {
 "mcpServerUrl": "http://risk-service:8091",
 "model": "claude-3-7-sonnet",
 "tools": [
],
 {
 "step": "gather_financial_data",
 "tool": "get_financial_statements"
 },
 {
 "step": "analyze_ratios",
 "tool": "calculate_financial_ratios"
 },
 {
 "step": "assess_industry_risk",
 "tool": "get_industry_trends"
 },
 {
 "step": "compute_final_score",
 "tool": "aggregate_risk_score"
 }
]
 },
 "mode": "SYNC_ENRICHMENT",
 "enrichmentOutputs": [
 "creditRiskScore",
 "riskCategory",
 "recommendedLimits",
 "mitigationSuggestions"
]
}
```

### Example Workflows

#### Workflow 1: Pure Rule-Based (Traditional DMN)

Use Case: Simple Product Configuration

No agents, just decision table evaluation.

```

{
 "templateId": "SIMPLE_PRODUCT_CONFIG",
 "entityType": "PRODUCT_CONFIGURATION",
 "agentConfig": {
 "enableAgents": false
 },
 "decisionTables": [
 {
 "name": "Product Approval Rules",
 "inputs": [
 {"name": "productType", "type": "string"},
 {"name": "priceVariance", "type": "number"}
],
 "outputs": [
 {"name": "approvalRequired", "type": "boolean"},
 {"name": "approverRoles", "type": "array"}
],
 "rules": [
 {
 "conditions": {"priceVariance": "<= 10"},
 "outputs": {"approvalRequired": false}
 },
 {
 "conditions": {"priceVariance": "> 10"},
 "outputs": {
 "approvalRequired": true,
 "approverRoles": ["PRODUCT_MANAGER"]
 }
 }
]
 }
]
}

```

#### Flow:

Submit → Validate → Evaluate DMN → Assign Approver (if needed) → Done

## Workflow 2: Async Agent Red Flag Detection

Use Case: Customer Onboarding with Fraud Detection

**Critical:** If agents detect fraud/AML issues, immediately terminate workflow.

```

{
 "templateId": "CUSTOMER_ONBOARDING_WITH_FRAUD_DETECTION",

```

```

"entityType": "CUSTOMER_ONBOARDING",
"agentConfig": {
 "enableAgents": true,
 "mode": "ASYNC_RED_FLAG",
 "agents": [
 {
 "agentId": "fraud-detection-mcp",
 "type": "MCP",
 "config": {
 "mcpServerUrl": "http://fraud-detection:8090",
 "model": "claude-3-7-sonnet",
 "tools": ["analyze_fraud_patterns", "check_device_risk"]
 },
 "redFlagConditions": {
 "fraudProbability": "> 0.75",
 "deviceRiskScore": "> 80",
 "identityMismatch": "true"
 },
 "redFlagAction": {
 "action": "TERMINATE_WORKFLOW",
 "autoReject": true,
 "reason": "High fraud risk detected by AI agent",
 "notifyRoles": ["FRAUD_TEAM", "COMPLIANCE_OFFICER"]
 }
 },
 {
 "agentId": "aml-sanctions-mcp",
 "type": "MCP",
 "config": {
 "mcpServerUrl": "http://aml-service:8092",
 "tools": ["check_sanctions_lists", "screen_pep"]
 },
 "redFlagConditions": {
 "sanctionsMatch": "true",
 "pepMatch": "true",
 "highRiskJurisdiction": "true"
 },
 "redFlagAction": {
 "action": "TERMINATE_WORKFLOW",
 "autoReject": true,
 "reason": "AML sanctions or PEP match found",
 "escalateTo": "AML_OFFICER"
 }
 }
],
 "timeout": 60000,
 "continueOnTimeout": false
},
"decisionTables": [
 {
 "name": "Onboarding Approval Rules",
 "note": "Only evaluated if NO red flags from agents",
 "inputs": [

```

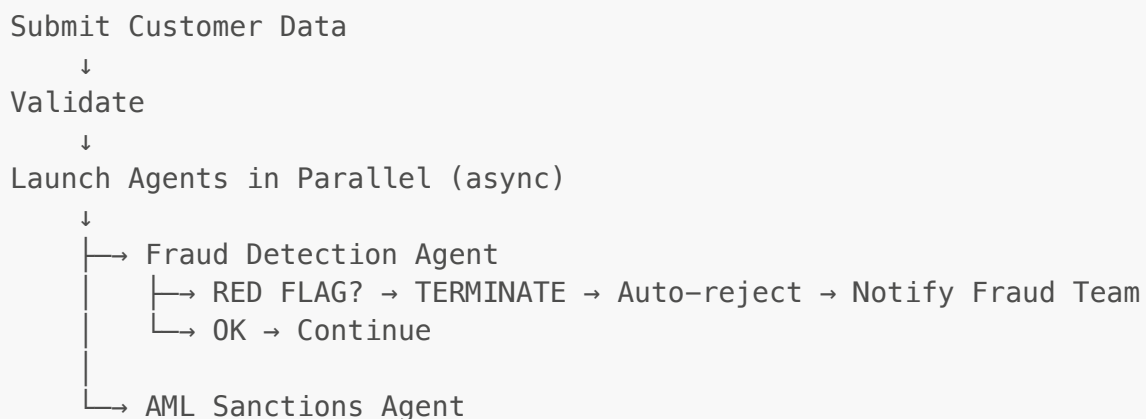


```

 {"name": "customerType", "type": "string"},
 {"name": "accountValue", "type": "number"},
 {"name": "agentFraudScore", "type": "number"},
 {"name": "agentAmlRisk", "type": "string"}
],
 "outputs": [
 {"name": "approverRoles", "type": "array"},
 {"name": "approvalCount", "type": "number"}
],
 "rules": [
 {
 "conditions": {
 "customerType": "INDIVIDUAL",
 "accountValue": "<= 100000",
 "agentFraudScore": "<= 30",
 "agentAmlRisk": "LOW"
 },
 "outputs": {
 "approverRoles": ["ONBOARDING_SPECIALIST"],
 "approvalCount": 1
 }
 },
 {
 "conditions": {
 "agentFraudScore": "> 30 && <= 60",
 "agentAmlRisk": "MEDIUM"
 },
 "outputs": {
 "approverRoles": ["ONBOARDING_SPECIALIST",
"COMPLIANCE_OFFICER"],
 "approvalCount": 2
 }
 }
]
}

```

### Flow:



```

 |→ RED FLAG? → TERMINATE → Auto-reject → Escalate to AML
Officer |→ OK → Continue
 ↓
Aggregate Agent Results
 ↓
Enrich entityMetadata with:
 - agentFraudScore: 25
 - agentAmlRisk: "LOW"
 ↓
Evaluate DMN with enriched data
 ↓
Result: Single approval needed (ONBOARDING_SPECIALIST)
 ↓
Human Approval

```

## Workflow 3: Sync GraphRAG + MCP + Final DMN

Use Case: Complex Loan Application with Knowledge Retrieval

**Pattern:** Agents run synchronously to enrich data, then DMN evaluates with agent insights.

```

{
 "templateId": "LOAN_APPLICATION_INTELLIGENT",
 "entityType": "LOAN_APPLICATION",
 "agentConfig": {
 "enableAgents": true,
 "mode": "SYNC_ENRICHMENT",
 "agents": [
 {
 "agentId": "credit-risk-graphrag",
 "type": "GRAPH_RAG",
 "config": {
 "graphDbUrl": "neo4j://credit-graph:7687",
 "embeddingModel": "text-embedding-3-large",
 "queries": [
 {
 "name": "get_customer_history",
 "cypher": "MATCH (c:Customer {id: $customerId})-[:HAS_LOAN]->(l:Loan) RETURN l.status, l.defaulted, l.amount"
 },
 {
 "name": "get_industry_risk",
 "cypher": "MATCH (i:Industry {name: $industry})-[:HAS_DEFAULT_RATE]->(r:RiskMetric) RETURN r.rate, r.trend"
 },
 {
 "name": "find_similar_customers",
 "cypher": "MATCH (c1:Customer {id: $customerId})-"

```

```
[[:SIMILAR_T0]->(c2:Customer)-[:HAS_LOAN]->(l:Loan) WHERE l.defaulted = true RETURN count(l) as similarDefaults"
```

```
 }
]
},
"enrichmentOutputs": [
 "historicalDefaultRate",
 "industryRiskTrend",
 "similarCustomerDefaults",
 "graphRiskScore"
]
},
{
 "agentId": "financial-analysis-mcp",
 "type": "MCP",
 "config": {
 "mcpServerUrl": "http://financial-analysis:8093",
 "model": "claude-3-7-sonnet",
 "tools": [
 "analyze_financial_statements",
 "calculate_dscr",
 "assess_collateral_value",
 "predict_cash_flow"
],
 "reasoningSteps": [
 {
 "step": "parse_financials",
 "tool": "analyze_financial_statements",
 "input": "$.entityData.financialStatements"
 },
 {
 "step": "compute_ratios",
 "tool": "calculate_dscr"
 },
 {
 "step": "evaluate_collateral",
 "tool": "assess_collateral_value",
 "input": "$.entityData.collateral"
 },
 {
 "step": "forecast_repayment",
 "tool": "predict_cash_flow"
 }
]
 },
 "enrichmentOutputs": [
 "debtServiceCoverageRatio",
 "collateralCoverageRatio",
 "cashFlowStability",
 "financialHealthScore"
]
},
{
```

```

 "agentId": "regulatory-compliance-graphrag",
 "type": "GRAPH_RAG",
 "config": {
 "graphDbUrl": "neo4j://compliance-graph:7687",
 "queries": [
 {
 "name": "check_lending_limits",
 "cypher": "MATCH (r:Regulation {jurisdiction:
$jurisdiction})-[:SETS_LIMIT]->(l:LendingLimit) WHERE l.loanType =
$loanType RETURN l.maxAmount, l.maxLTV"
 },
 {
 "name": "verify_documentation",
 "cypher": "MATCH (r:Regulation)-[:REQUIRES_DOCUMENT]->
(d:DocumentRequirement) WHERE r.jurisdiction = $jurisdiction RETURN
d.documentType, d.mandatory"
 }
]
 },
 "enrichmentOutputs": [
 "regulatoryMaxLoanAmount",
 "maxLoanToValue",
 "requiredDocuments",
 "complianceStatus"
]
 },
 "executionMode": "SEQUENTIAL",
 "failOnAgentError": false,
 "timeout": 120000
},
"decisionTables": [
 {
 "name": "Loan Approval Decision",
 "note": "Evaluates using both original data AND agent-enriched
insights",
 "inputs": [
 {"name": "loanAmount", "type": "number"},
 {"name": "creditScore", "type": "number"},
 {"name": "loanType", "type": "string"},
 {"name": "graphRiskScore", "type": "number", "source": "agent"},
 {"name": "debtServiceCoverageRatio", "type": "number", "source":
"agent"},
 {"name": "collateralCoverageRatio", "type": "number", "source":
"agent"},
 {"name": "historicalDefaultRate", "type": "number", "source":
"agent"},
 {"name": "complianceStatus", "type": "string", "source":
"agent"}
],
 "outputs": [
 {"name": "decision", "type": "string"},
 {"name": "approverRoles", "type": "array"},

```

```

 {"name": "approvalCount", "type": "number"},
 {"name": "conditions", "type": "array"}
],
 "rules": [
 {
 "ruleId": "AUTO_REJECT_COMPLIANCE_FAIL",
 "priority": 100,
 "conditions": {
 "complianceStatus": "FAILED"
 },
 "outputs": {
 "decision": "AUTO_REJECT",
 "approverRoles": [],
 "conditions": ["Regulatory compliance requirements not met"]
 }
 },
 {
 "ruleId": "AUTO_REJECT_HIGH_RISK",
 "priority": 99,
 "conditions": {
 "graphRiskScore": "> 80",
 "historicalDefaultRate": "> 15"
 },
 "outputs": {
 "decision": "AUTO_REJECT",
 "approverRoles": [],
 "conditions": ["High risk based on historical patterns and
graph analysis"]
 }
 },
 {
 "ruleId": "AUTO_APPROVE_EXCELLENT",
 "priority": 95,
 "conditions": {
 "creditScore": ">= 750",
 "debtServiceCoverageRatio": ">= 1.5",
 "collateralCoverageRatio": ">= 1.3",
 "graphRiskScore": "<= 30",
 "loanAmount": "<= 500000"
 },
 "outputs": {
 "decision": "AUTO_APPROVE",
 "approverRoles": [],
 "approvalCount": 0,
 "conditions": ["Excellent credit metrics and low AI-assessed
risk"]
 }
 },
 {
 "ruleId": "SINGLE_APPROVAL_GOOD",
 "priority": 80,
 "conditions": {
 "creditScore": ">= 680",

```

```

 "debtServiceCoverageRatio": ">= 1.25",
 "graphRiskScore": "<= 50"
 },
 "outputs": {
 "decision": "APPROVE_WITH_CONDITIONS",
 "approverRoles": ["CREDIT_OFFICER"],
 "approvalCount": 1,
 "conditions": ["Standard approval required"]
 }
},
{
 "ruleId": "DUAL_APPROVAL_MODERATE_RISK",
 "priority": 75,
 "conditions": {
 "creditScore": ">= 620",
 "graphRiskScore": "> 50 && <= 70",
 "debtServiceCoverageRatio": ">= 1.15"
 },
 "outputs": {
 "decision": "APPROVE_WITH_CONDITIONS",
 "approverRoles": ["CREDIT_OFFICER",
"SENIOR_CREDIT_OFFICER"],
 "approvalCount": 2,
 "conditions": [
 "Enhanced monitoring required",
 "Consider additional collateral"
]
 }
},
{
 "ruleId": "COMMITTEE_APPROVAL_HIGH_VALUE",
 "priority": 90,
 "conditions": {
 "loanAmount": "> 5000000"
 },
 "outputs": {
 "decision": "COMMITTEE_APPROVAL",
 "approverRoles": ["CREDIT_COMMITTEE"],
 "approvalCount": 1,
 "conditions": [
 "Credit committee review required for high-value loans"
]
 }
}
],
"hitPolicy": "PRIORITY"
}
]
}

```

**Flow:**

Submit Loan Application

↓

Validate

↓

Execute Agents Sequentially:

↓

1. Credit Risk GraphRAG Agent

- Query graph for customer history
- Query graph for industry risk
- Find similar customer patterns
- Output: graphRiskScore: 45, historicalDefaultRate: 8%, similarCustomerDefaults: 3

↓

2. Financial Analysis MCP Agent

- Step 1: Parse financial statements
- Step 2: Calculate DSCR = 1.4
- Step 3: Assess collateral coverage = 1.5
- Step 4: Predict cash flow stability = "HIGH"
- Output: debtServiceCoverageRatio: 1.4, collateralCoverageRatio: 1.5

↓

3. Regulatory Compliance GraphRAG Agent

- Check lending limits for jurisdiction
- Verify required documents
- Output: complianceStatus: "PASSED", maxLoanToValue: 80%

↓

Aggregate Agent Results & Enrich Metadata:

```
{
 loanAmount: 750000,
 creditScore: 720,
 loanType: "COMMERCIAL",
 graphRiskScore: 45, ← from GraphRAG
 debtServiceCoverageRatio: 1.4, ← from MCP
 collateralCoverageRatio: 1.5, ← from MCP
 historicalDefaultRate: 8, ← from GraphRAG
 complianceStatus: "PASSED" ← from GraphRAG
}
```

↓

Evaluate DMN with Enriched Data:

- Rule AUTO\_REJECT\_COMPLIANCE\_FAIL: NO
- Rule AUTO\_REJECT\_HIGH\_RISK: NO (graphRiskScore=45 < 80)
- Rule AUTO\_APPROVE\_EXCELLENT: NO (loanAmount=750k > 500k)
- Rule SINGLE\_APPROVAL\_GOOD: YES ✓

↓

Result:

```
{
 decision: "APPROVE_WITH_CONDITIONS",
 approverRoles: ["CREDIT_OFFICER"],
 approvalCount: 1,
 conditions: ["Standard approval required"]
}
```

↓

Assign CREDIT\_OFFICER

## Agent Decision Models

### AgentDecision

```
@Data
@Builder
public class AgentDecision {
 private String agentId;
 private String agentType; // "MCP", "GRAPH_RAG"
 private LocalDateTime executedAt;
 private Duration executionTime;

 // Red flag detection
 private boolean redFlagDetected;
 private String redFlagReason;
 private RedFlagSeverity severity;
 private RedFlagAction recommendedAction;

 // Enrichment outputs
 private Map<String, Object> enrichmentData;

 // Agent reasoning trace
 private List<AgentReasoningStep> reasoningSteps;

 // Confidence and metadata
 private double confidenceScore;
 private String model;
 private Map<String, Object> agentMetadata;
}

public enum RedFlagSeverity {
 LOW,
 MEDIUM,
 HIGH,
 CRITICAL
}

public enum RedFlagAction {
 CONTINUE, // Log but continue
 ENHANCE_REVIEW, // Add additional approver
 TERMINATE_REJECT, // Auto-reject
 ESCALATE // Escalate to senior
}
```

### AgentReasoningStep



```

@Data
@Builder
public class AgentReasoningStep {
 private int stepNumber;
 private String stepName;
 private String tool;
 private Map<String, Object> input;
 private Map<String, Object> output;
 private String reasoning;
 private LocalDateTime timestamp;
 private Duration duration;
}

```

## MCPAgentConfig

```

@Data
@Builder
public class MCPAgentConfig {
 private String mcpServerUrl;
 private String model;
 private List<String> tools;
 private List<ReasoningStep> reasoningSteps;
 private int timeout;
 private Map<String, String> headers;
 private RetryPolicy retryPolicy;
}

@Data
public class ReasoningStep {
 private String step;
 private String tool;
 private String input; // JSONPath expression
 private Map<String, Object> parameters;
}

```

## GraphRAGAgentConfig

```

@Data
@Builder
public class GraphRAGAgentConfig {
 private String graphDbUrl;
 private String embeddingModel;
 private int retrievalDepth;
 private List<String> entities;
 private List<CypherQuery> queries;
 private SemanticSearchConfig semanticSearch;
}

```

```

@Data
public class CypherQuery {
 private String name;
 private String cypher;
 private Map<String, String> parameters;
 private String resultMapping;
}

@Data
public class SemanticSearchConfig {
 private boolean enabled;
 private String questionField;
 private int topK;
 private double similarityThreshold;
}

```

## Agent Execution Models

### AgentExecutionPlan

```

@Data
@Builder
public class AgentExecutionPlan {
 private ExecutionMode mode;
 private List<AgentTask> tasks;
 private int timeout;
 private boolean failOnError;
 private AgentOrchestrationStrategy strategy;
}

public enum ExecutionMode {
 ASYNC_RED_FLAG, // Parallel, terminate on red flag
 SYNC_ENRICHMENT, // Sequential, enrich metadata
 HYBRID // Mix of both
}

public enum AgentOrchestrationStrategy {
 PARALLEL, // All agents execute simultaneously
 SEQUENTIAL, // One after another
 CONDITIONAL, // Based on previous results
 PRIORITY_BASED // High priority first
}

```

### AgentTask

```

@Data
@Builder

```

```

public class AgentTask {
 private String agentId;
 private AgentType type;
 private Object config; // MCPAgentConfig or GraphRAGAgentConfig
 private ExecutionMode mode;

 // Red flag configuration
 private Map<String, String> redFlagConditions;
 private RedFlagAction redFlagAction;

 // Enrichment configuration
 private List<String> enrichmentOutputs;
 private String outputMapping;

 // Execution settings
 private int priority;
 private int timeout;
 private RetryPolicy retryPolicy;
}

public enum AgentType {
 MCP,
 GRAPH_RAG,
 CUSTOM
}

```

## Async Red Flag Pattern

### Workflow Behavior

```

// In Temporal workflow
@WorkflowMethod
public WorkflowResult execute(WorkflowSubject subject) {

 // Step 1: Launch all agents
 List<Promise<AgentDecision>> agentPromises = new ArrayList<>();

 for (AgentTask task : subject.getAgentConfig().getTasks()) {
 Promise<AgentDecision> promise = Async.function(() ->
 executeAgent(task, subject)
);
 agentPromises.add(promise);
 }

 // Step 2: Wait for first red flag OR all complete
 Selector selector = new Selector();

 for (Promise<AgentDecision> promise : agentPromises) {
 selector.addCondition(() -> promise.isCompleted() &&
 promise.get().isRedFlagDetected(),

```

```

 () -> {
 AgentDecision decision = promise.get();
 // RED FLAG DETECTED – TERMINATE
 return WorkflowResult.builder()
 .success(false)
 .resultCode("RED_FLAG_DETECTED")
 .message(decision.getRedFlagReason())
 .agentDecision(decision)
 .build();
 });
 }

 // All agents complete without red flags
 selector.addCondition(() ->
 Promise.allOf(agentPromises).isCompleted(),
 () -> {
 // Aggregate agent results
 List<AgentDecision> decisions = agentPromises.stream()
 .map(Promise::get)
 .collect(Collectors.toList());

 // Enrich metadata
 Map<String, Object> enrichedMetadata = enrichMetadata(
 subject.getEntityMetadata(),
 decisions
);

 // Evaluate DMN with enriched data
 ComputedApprovalPlan plan = evaluateRules(
 template,
 enrichedMetadata
);

 return continueWorkflow(plan);
 });

 selector.select();
}

```

## Integration with MCP Servers

### MCP Tool Invocation

```

public class MCPAgentExecutor {

 public AgentDecision execute(MCPAgentConfig config, WorkflowSubject
subject) {

 List<AgentReasoningStep> steps = new ArrayList<>();
 Map<String, Object> context = new HashMap<>

```

```

(subject.getEntityData());

// Execute reasoning steps sequentially
for (ReasoningStep step : config.getReasoningSteps()) {

 // Invoke MCP tool
 MCPToolRequest request = MCPToolRequest.builder()
 .tool(step.getTool())
 .input(extractInput(step.getInput(), context))
 .parameters(step.getParameters())
 .build();

 MCPToolResponse response = mcpClient.invokeTool(
 config.getMcpServerUrl(),
 request
);

 // Record reasoning step
 steps.add(AgentReasoningStep.builder()
 .stepNumber(steps.size() + 1)
 .stepName(step.getStep())
 .tool(step.getTool())
 .input(request.getInput())
 .output(response.getResult())
 .reasoning(response.getReasoning())
 .build());

 // Add to context for next step
 context.putAll(response.getResult());
}

// Extract enrichment data
Map<String, Object> enrichment = extractEnrichment(
 context,
 config.getEnrichmentOutputs()
);

// Check for red flags
boolean redFlag = checkRedFlags(enrichment,
config.getRedFlagConditions());

return AgentDecision.builder()
 .agentId(config.getAgentId())
 .agentType("MCP")
 .redFlagDetected(redFlag)
 .enrichmentData(enrichment)
 .reasoningSteps(steps)
 .model(config.getModel())
 .build();
}
}

```

# Benefits

## 1. Intelligent Automation

- AI agents pre-screen applications
- Auto-reject obvious fraud/risk cases
- Auto-approve low-risk cases

## 2. Context-Aware Decisions

- GraphRAG retrieves relevant historical patterns
- Knowledge graph provides regulatory context
- MCP agents apply complex reasoning

## 3. Flexible Execution

- Async red flag detection for early termination
- Sync enrichment for rule enhancement
- Hybrid approaches for complex scenarios

## 4. Explainable AI

- Full reasoning trace from agents
- DMN rules remain interpretable
- Combined AI + human judgment

## 5. Extensible Agent Framework

- Add new MCP tools without code changes
- Configure GraphRAG queries via templates
- Plugin architecture for custom agents

## Configuration Example

```
application.yml
workflow:
 agents:
 mcp:
 default-timeout: 30000
 retry-attempts: 2
 servers:
 fraud-detection: http://fraud-service:8090
 financial-analysis: http://financial-service:8093
 aml-screening: http://aml-service:8092

 graph-rag:
 neo4j-url: neo4j://graph-db:7687
 embedding-service: http://embedding-service:8094
 default-retrieval-depth: 3
```

```
cache-ttl: 3600
```

```
orchestration:
```

```
max-parallel-agents: 5
```

```
default-timeout: 120000
```

```
fail-fast: true
```

This hybrid approach combines the best of rule-based and AI-driven decision-making!